

Exercises03-Solutions

February 21, 2015

```
In [1]: import os
import sys
import glob
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
%matplotlib inline
plt.style.use('ggplot')
%precision 4
```

```
Out[1]: u'%.4f'
```

Exercise 1 (10 pts). Avoiding catastrophic cancellation.

The tail of the standard logistic distribution is given by $1 - F(t) = 1 - (1 + e^{-t})^{-1}$.

- Define a function `f1` to calculate the tail probability of the logistic distribution using the formula given above
- Use `sympy` to find the exact value of the tail distribution (using the same symbolic formula) to 20 decimal digits
- Calculate the *relative error* of `f1` when $t = 25$ (The relative error is given by `abs(exact - approximate)/exact`)
- Rewrite the expression for the tail of the logistic distribution using simple algebra so that there is no risk of cancellation, and write a function `f2` using this formula. Calculate the *relative error* of `f2` when $t = 25$.
- How much more accurate is `f2` compared with `f1` in terms of the relative error?

```
In [2]: # Your code here
```

```
def f1(t):
    """Calculates tail probability of the logistic distribution."""
    return 1 - 1.0/(1 + np.exp(-t))

def fsymb(t, n=20):
    """Exact value to n decimal digits using symbolic algebra."""
    from sympy import exp
    return (1 - 1/(1 + exp(-t))).evalf(n=n)

def f2(t):
    """Calculates tail probability of the logistic distribution - no cancellation."""
    return 1/(1 + np.exp(t))

r1 = abs(fsymb(25) - f1(25))/fsymb(25)
r2 = abs(fsymb(25) - f2(25))/fsymb(25)
```

```

print "Relative error of f1:\t%.16f" % r1
print "Relative error of f2\t%.16f" % r2
print "f2 improvieemnt over f1\t%g" % (r1/r2)

```

```

Relative error of f1:          0.0000041759147666
Relative error of f2          0.000000000000000001
f2 improvieemnt over f1      3.66247e+10

```

Exercise 2 (10 pts). Ill-conditioned linear problems.

You are given a $n \times p$ design matrix X and a p -vector of observations y and asked to find the coefficients β that solve the linear equations $X\beta = y$.

```

X = np.load('x.npy')
y = np.load('y.npy')

```

The solution β can also be loaded as

```
beta = np.load('b.npy')
```

- Write a formulat that could solve the system of linear equations in terms of X and y Write a function `f1` that takes arguments X and y and returns β using this.
- How could you code this formula using `np.linalg.solve` tht does not require inverting a matrix? Write a function `f2` that takes arguments X and y and returns β using this.
- Note that carefully designed algorithms *can* solve this ill-conditioned problem, which is why you should always use library functions for linear algebra rather than write your own.

```
np.linalg.lstsq(x, y)[0]
```

- What happens if you try to solve for β using `f1` or `f2`? Remove the column of X that is making the matrix singular and find the $p - 1$ vector b using `f2`.
- Note that the solution differs from that given by `np.linalg.lstsq`. This arises because the relevant condition number for `f2` is actually for the matrix $X^T X$ while the condition number of `lstsq` is for the matrix X . Why is the condition so high even after removing the column that makes the matrix singular?

```

In [3]: X = np.load('x.npy')
        y = np.load('y.npy')
        beta = np.load('b.npy')

```

```
# Your code here
```

```

def f1(X, y):
    """Direct translation of normal equations to code."""
    return np.dot(np.linalg.inv(np.dot(X.T, X)), np.dot(X.T, y))

```

```

def f2(X, y):
    """Solving normal equations wihtout matrix inversion."""
    return np.solve(np.dot(x.T, x), np.dot(x.T, np.dot(x, b)))

```

```

%precision 2
print "X = "
print X
# counting from 0 (so column 5 is the last column)
# we can see that column 5 is a multiple of column 3

```

```

# so one approach is to simply remove this (dependent) column

print "True solution\t\t", beta
print "Library function\t", np.linalg.lstsq(X, y)[0]
print "Using f2\t\t", f1(X[:, :5], y)

# Condition number is still high because column 1 is on a much
# larger scale than all the other columns

X =
[[ 5.00e+00  4.82e+14  9.00e+00  5.00e+00  0.00e+00  5.00e+01]
 [ 1.00e+00  4.21e+14  6.00e+00  9.00e+00  2.00e+00  9.00e+01]
 [ 5.00e+00  1.20e+14  4.00e+00  2.00e+00  4.00e+00  2.00e+01]
 [ 7.00e+00  5.42e+14  1.00e+00  7.00e+00  0.00e+00  7.00e+01]
 [ 9.00e+00  5.42e+14  7.00e+00  6.00e+00  9.00e+00  6.00e+01]
 [ 0.00e+00  6.02e+13  8.00e+00  8.00e+00  3.00e+00  8.00e+01]
 [ 8.00e+00  4.21e+14  3.00e+00  6.00e+00  5.00e+00  6.00e+01]
 [ 9.00e+00  1.81e+14  4.00e+00  8.00e+00  1.00e+00  8.00e+01]
 [ 0.00e+00  1.81e+14  9.00e+00  2.00e+00  0.00e+00  2.00e+01]
 [ 9.00e+00  1.20e+14  7.00e+00  7.00e+00  9.00e+00  7.00e+01]]
True solution          [ 0.47  0.1  0.9  0.12  0.52  0.08]
Library function       [ 0.47  0.1  0.9  0.01  0.52  0.09]
Using f2               [ 0.47  0.1  0.9  0.95  0.53]

```

Exercise 3 (10 pts). Importance of using efficient algorithms.

- Implement bubble sort
- Calculate its big \mathcal{O} algorithmic complexity
- Time the performance of bubble sort on random uniform deviate vectors of sizes `range(100, 2000, 100)` using `time.time()` from the standard library
- Use `scipy.optimize.curve_fit` to fit an appropriate function to the collection of (size, execution time) data points. Extrapolate how long it would take to sort a random vector of size 1,000,000. Now time how long it takes for the system sort to sort a random vector of size 1,000,000.
- Plot the fits together with the data points using `matplotlib.pyplot` functions.

In [24]: # Your code here

```

def bubble(xs):
    """Bubble sort."""
    for i in range(len(xs)):
        for j in range(i, len(xs)):
            if xs[i] > xs[j]:
                xs[i], xs[j] = xs[j], xs[i]
    return xs

import time

ns = range(100, 2000, 100)
bubble_times = []

for n in ns:
    xs = np.random.random(n)
    start = time.time()
    bubble(xs)
    bubble_times.append(time.time() - start)

```

```

def func(x, a, b, c):
    """Quadratic function for O(n^2) complexity."""
    return a*x**2 + b*x + c

from scipy.optimize import curve_fit

plt.scatter(ns, bubble_times, c='blue')
popt, pcov = curve_fit(func, ns, bubble_times)
a, b, c = popt

print "Predicted time to sort 1,000,000 items = %.2f seconds" % func(1000000, a, b, c)


x = np.random.random(1000000)
start = time.time()
x.sort()
elapsed = time.time() - start
print "Time for system sort to sort 1,000,000 items = %.2f seconds" % elapsed

xp = np.linspace(0, 2000)
plt.plot(xp, func(xp, a, b, c))
plt.xlim([0, 2000]);
plt.ylim([0, plt.ylim()[1]])
plt.xlabel('Length of vector n')
plt.ylabel('Time to sort (s)');

```

Predicted time to sort 1,000,000 items = 290278.71 seconds
Time for system sort to sort 1,000,000 items = 0.14 seconds

Out[24]: <matplotlib.text.Text at 0x115cd9750>



Exercise 4 (20 pts). One of the goals of the course is that you will be able to implement novel algorithms from the literature.

- Implement the mean-shift algorithm in 1D as described [here](#).
 - Use the following function signature


```
def mean_shift(xs, m, kernel, max_iters=100, tol=1e-6):
```
 - `xs` is the data set, `m` is the starting location, and `kernel` is a kernel function
 - `tol` is the difference in $\|x\|$ across iterations

- Use the following kernels with bandwidth h (a default value of 1.0 will work fine)
 - Flat - return 1 if $\|x\| < h$ and 0 otherwise
 - Gaussian

$$\frac{1}{\sqrt{2\pi}h} e^{-\frac{\|x\|^2}{h^2}}$$

- Note that $\|x\|$ is the norm of the data point being evaluated minus the current value of x
- Use both kernels to find all 3 modes of the data set in `x1d.npy`
- Modify the algorithm and/or kernels so that it now works in an arbitrary number of dimensions.
- Use both kernels to find all 3 modes of the data set in `x2d.npy`
- Plot the path of successive intermediate solutions of the mean-shift algorithm starting from `x0 = (-4, 10)` until it converges onto a mode in the 2D data for each kernel. Superimpose the path on top of a contour plot of the data density.

In [101]: *# Your code here*

```
def gaussian_kernel(xs, x, h=1.0):
    """Gaussian kernel for a shifting window centered at x."""
    X = xs - x
    try:
        d = xs.shape[1]
    except:
        d = 1
    k = np.array([(2*np.pi*h**d)**-0.5*np.exp(-(np.dot(_.T, _)/h)**2) for _ in X])
    if d != 1:
        k = k[:, np.newaxis]
    return k

def flat_kernel(xs, x, h=1.0):
    """Flat kernel for a shifting window centered at x."""
    X = xs - x
    try:
        d = xs.shape[1]
    except:
        d = 1
    k = np.array([1 if np.dot(_.T, _) < h else 0 for _ in X])
    if d != 1:
        k = k[:, np.newaxis]
    return k
```

```

def mean_shift(xs, x, kernel, max_iters=100, tol=1e-6, trace=False):
    """Finds the local mode using mean shift algorithm."""

    record = []

    for i in range(max_iters):
        if trace:
            record.append(x)
        m = (kernel(xs, x)*xs).sum(axis=0)/kernel(xs, x).sum(axis=0) - x
        if np.sum(m**2) < tol:
            break
        x += m
    return i, x, np.array(record)

In [110]: x1 = np.load('x1d.npy')

# choose kernel to evaluate
kernel = flat_kernel
# kernel = gaussian_kernel

i1, m1, path = mean_shift(x1, 1, kernel)
print i1, m1


i2, m2, path = mean_shift(x1, -7, kernel)
print i2, m2

i3, m3, path = mean_shift(x1, 7, kernel)
print i3, m3

xp = np.linspace(0, 1.0, 100)
plt.hist(x1, 50, histtype='step', normed=True);
plt.axvline(m1, c='blue')
plt.axvline(m2, c='blue')
plt.axvline(m3, c='blue');

15 0.575467435455
16 -4.99502467965
16 3.97936417657

```



```
In [123]: x2 = np.load('x2d.npy')

# choose kernel to evaluate
# kernel = flat_kernel
kernel = gaussian_kernel

i1, m1, path1 = mean_shift(x2, [0,0], kernel, trace=True)
print i1, m1

i2, m2, path2 = mean_shift(x2, [-4,5], kernel, trace=True)
print i2, m2

i3, m3, path3 = mean_shift(x2, [10,10], kernel, trace=True)
print i3, m3

59 [ 2.318177222  2.825508637]
12 [-3.0704798654  3.0567861169]
42 [ 6.0225434115  8.9506699557]

In [155]: import scipy.stats as stats

# size of marekr at starting position
base = 40

# set up for estimating density using gaussian_kde
xmin, xmax = -6, 12
ymin, ymax = -5, 15
X, Y = np.mgrid[xmin:xmax:100j, ymin:ymax:100j]
positions = np.vstack([X.ravel(), Y.ravel()])
```

```

kde = stats.gaussian_kde(x2.T)
Z = np.reshape(kde(positions).T, X.shape)


plt.contour(X, Y, Z)
# plot data in background
plt.scatter(x2[:, 0], x2[:, 1], c='grey', alpha=0.2, edgecolors='none')

# path from [0,0]
plt.scatter(path1[:, 0], path1[:, 1], s=np.arange(base, base+len(path1)),
            c='none', edgecolors='red', marker='x', linewidth=1.5)

# path from [-4,5]
plt.scatter(path2[:, 0], path2[:, 1], s=np.arange(base, base+len(path2)),
            c='none', edgecolors='blue', marker='x', linewidth=1.5)

# path from [10,10]
plt.scatter(path3[:, 0], path3[:, 1], s=np.arange(base, base+len(path3)),
            c='none', edgecolors='green', marker='x', linewidth=1.5)
plt.axis([xmin, xmax, ymin, ymax]);

```



In []: