

UsingPandas

February 21, 2015

```
In [1]: %matplotlib inline
import numpy as np
import matplotlib.pyplot as plt

In [2]: import warnings
warnings.filterwarnings("ignore")
```

0.1 Using Pandas

The `numpy` module is excellent for numerical computations, but to handle missing data or arrays with mixed types takes more work. The `pandas` module provides objects similar to R's data frames, and these are more convenient for most statistical analysis. The `pandas` module also provides many methods for data import and manipulation that we will explore in this section.

[Pandas for R Users](#)

```
In [3]: import pandas as pd
import statsmodels.api as sm
from pandas import Series, DataFrame, Panel
from string import ascii_lowercase as letters
from scipy.stats import chisqprob
```

0.1.1 Series

Series is a 1D array with axis labels.

```
In [4]: # Creating a series and extracting elements.

xs = Series(np.arange(10), index=tuple(letters[:10]))
print xs[:3], '\n'
print xs[7:], '\n'
print xs[:3], '\n'
print xs[['d', 'f', 'h']], '\n'
print xs.d, xs.f, xs.h

a    0
b    1
c    2
dtype: int64

h    7
i    8
j    9
dtype: int64
```

```
a    0
d    3
g    6
j    9
dtype: int64
```

```
d    3
f    5
h    7
dtype: int64
```


```
3 5 7
```

In [5]: *# All the numpy functions will work with Series objects, and return another Series*

```
y1, y2 = np.mean(xs), np.var(xs)
y1, y2
```

Out[5]: (4.5, 8.25)

In [6]: *# Matplotlib will work on Series objects too*
plt.plot(xs, np.sin(xs), 'r-o', xs, np.cos(xs), 'b-x');



In [7]: *# Convert to numpy arrays with values*

```
print xs.values
```

```
[0 1 2 3 4 5 6 7 8 9]
```

In [8]: *# The Series datatype can also be used to represent time series*

```

import datetime as dt
from pandas import date_range

# today = dt.date.today()
today = dt.datetime.strptime('Jan 21 2015', '%b %d %Y')
print today, '\n'
days = date_range(today, periods=35, freq='D')
ts = Series(np.random.normal(10, 1, len(days)), index=days)

# Extracting elements
print ts[0:4], '\n'
print ts['2015-01-21':'2015-01-28'], '\n' # Note - includes end time

```

```
2015-01-21 00:00:00
```

```

2015-01-21    9.719261
2015-01-22    8.894461
2015-01-23   10.074521
2015-01-24   10.769334
Freq: D, dtype: float64

```

```

2015-01-21    9.719261
2015-01-22    8.894461
2015-01-23   10.074521
2015-01-24   10.769334
2015-01-25   10.159401
2015-01-26    8.992754
2015-01-27    9.681121
2015-01-28    9.908445
Freq: D, dtype: float64

```

In [9]: *# We can generate statistics for time ranges with the resample method*
For example, suppose we are interested in weekly means, standard deviations and sum-of-square

```

df = ts.resample(rule='W', how=('mean', 'std', lambda x: sum(x*x)))
df

```

```

Out[9]:
          mean      std  <lambda>
2015-01-25  9.923396  0.688209  494.263430
2015-02-01 10.357088  0.848930  755.208973
2015-02-08 10.224806  0.869441  736.362134
2015-02-15 10.672230  0.942680  802.607338
2015-02-22  9.785174  1.012906  676.403270
2015-03-01  9.495084  1.472653  182.481942

```

0.1.2 DataFrame

For statisticians, a DataFrame is similar to the R dataframe object. For everyone else, it is like a simple tabular spreadsheet. Each column is a Series object.

In [10]: *# Note that the df object in the previous cell is a DataFrame*
print type(df)

```
<class 'pandas.core.frame.DataFrame'>
```

```
In [11]: # Renaming columns
# The use of mean and std are problematic because there are also methods in dataframe with tho
# Also, <lambda> is uninformative
# So we would like to give better names to the columns of df
```

```
df.columns = ('mu', 'sigma', 'sum_of_sq')
print df
```

mu	sigma	sum_of_sq
2015-01-25	9.923396	0.688209 494.263430
2015-02-01	10.357088	0.848930 755.208973
2015-02-08	10.224806	0.869441 736.362134
2015-02-15	10.672230	0.942680 802.607338
2015-02-22	9.785174	1.012906 676.403270
2015-03-01	9.495084	1.472653 182.481942

```
In [12]: # Extracting columns from a DataFrame
```

```
print df.mu, '\n' # by attribute
print df['sigma'], '\n' # by column name
```

```
2015-01-25      9.923396
2015-02-01     10.357088
2015-02-08     10.224806
2015-02-15     10.672230
2015-02-22      9.785174
2015-03-01      9.495084
Freq: W-SUN, Name: mu, dtype: float64
```

```
2015-01-25      0.688209
2015-02-01      0.848930
2015-02-08      0.869441
2015-02-15      0.942680
2015-02-22      1.012906
2015-03-01      1.472653
Freq: W-SUN, Name: sigma, dtype: float64
```

```
In [13]: # Extracting rows from a DataFrame
```

```
print df[1:3], '\n'
print df['2015-01-21'::2]
```

mu	sigma	sum_of_sq
2015-02-01	10.357088	0.848930 755.208973
2015-02-08	10.224806	0.869441 736.362134

	mu	sigma	sum_of_sq
2015-01-25	9.923396	0.688209	494.263430
2015-02-08	10.224806	0.869441	736.362134
2015-02-22	9.785174	1.012906	676.403270

```
In [14]: # Extracting blocks and scalars
```

```
print df.iat[2, 2], '\n' # extract an element with iat()
print df.loc['2015-01-25':'2015-03-01', 'sum_of_sq'], '\n' # indexing by label
print df.iloc[:3, 2], '\n' # indexing by position
print df.ix[:3, 'sum_of_sq'], '\n' # by label OR position
```

736.362134378

```
2015-01-25    494.263430
2015-02-01    755.208973
2015-02-08    736.362134
2015-02-15    802.607338
2015-02-22    676.403270
2015-03-01    182.481942
```

Freq: W-SUN, Name: sum_of_sq, dtype: float64

```
2015-01-25    494.263430
2015-02-01    755.208973
2015-02-08    736.362134
```

Freq: W-SUN, Name: sum_of_sq, dtype: float64

```
2015-01-25    494.263430
2015-02-01    755.208973
2015-02-08    736.362134
```

Freq: W-SUN, Name: sum_of_sq, dtype: float64

In [15]: *# Using Boolean conditions for selecting elements*

```
print df[(df.sigma < 1) & (df.sum_of_sq < 700)], '\n' # need parenthesis because of operator p
print df.query('sigma < 1 and sum_of_sq < 700') # the query() method allows more readable quer
```

```
mu      sigma  sum_of_sq
2015-01-25  9.923396  0.688209  494.26343
```

```
mu      sigma  sum_of_sq
2015-01-25  9.923396  0.688209  494.26343
```

0.1.3 Panels

Panels are 3D arrays - they can be thought of as dictionaries of DataFrames.

```
In [16]: df= np.random.binomial(100, 0.95, (9,2))
dm = np.random.binomial(100, 0.9, [12,2])
dff = DataFrame(df, columns = ['Physics', 'Math'])
dfm = DataFrame(dm, columns = ['Physics', 'Math'])
score_panel = Panel({'Girls': dff, 'Boys': dfm})
print score_panel, '\n'
```

```
<class 'pandas.core.panel.Panel'>
Dimensions: 2 (items) x 12 (major_axis) x 2 (minor_axis)
Items axis: Boys to Girls
Major_axis axis: 0 to 11
Minor_axis axis: Physics to Math
```

```
In [17]: score_panel['Girls'].transpose()
```

```
Out[17]:
```

	0	1	2	3	4	5	6	7	8	9	10	11
Physics	95	95	96	95	93	95	96	94	96	NaN	NaN	NaN
Math	95	95	94	92	91	92	96	95	97	NaN	NaN	NaN

```
In [18]: # find physics and math scores of girls who scored >= 93 in math
# a DataFrame is returned
score_panel.ix['Girls', score_panel.Girls.Math >= 93, :]
```

```
Out[18]:
```

	Physics	Math
0	95	95
1	95	95
2	96	94
6	96	96
7	94	95
8	96	97

0.1.4 Split-Apply-Combine

Many statistical summaries are in the form of split along some property, then apply a function to each subgroup and finally combine the results into some object. This is known as the ‘split-apply-combine’ pattern and implemented in Pandas via `groupby()` and a function that can be applied to each subgroup.

```
In [19]: # import a DataFrame to play with
try:
    tips = pd.read_pickle('tips.pic')
except:
    tips = pd.read_csv('https://raw.githubusercontent.com/vincentarelbundock/Rdatasets/master/csv/reshape2/tips.csv')
    tips.to_pickle('tips.pic')
```

```
In [20]: tips.head(n=4)
```

```
Out[20]:
```

	Unnamed: 0	total_bill	tip	sex	smoker	day	time	size
0	1	16.99	1.01	Female	No	Sun	Dinner	2
1	2	10.34	1.66	Male	No	Sun	Dinner	3
2	3	21.01	3.50	Male	No	Sun	Dinner	3
3	4	23.68	3.31	Male	No	Sun	Dinner	2

```
In [21]: # We have an extra set of indices in the first column
# Let's get rid of it
```

```
tips = tips.ix[:, 1:]
tips.head(n=4)
```

```
Out[21]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2

```
In [22]: # For an example of the split-apply-combine pattern, we want to see counts by sex and smoker status
# In other words, we split by sex and smoker status to get 2x2 groups,
# then apply the size function to count the number of entries per group
# and finally combine the results into a new multi-index Series
```

```
grouped = tips.groupby(['sex', 'smoker'])
grouped.size()
```

```
Out[22]:
```

sex	smoker	
Female	No	54
	Yes	33
Male	No	97
	Yes	60

dtype: int64

```
In [37]: # If you need the margins, use the crosstab function
```

```
pd.crosstab(tips.sex, tips.smoker, margins=True)
```

```
Out[37]: smoker    No  Yes  All
sex
Female    54   33   87
Male     97   60  157
All     151   93  244
```

```
In [23]: # If more than 1 column of results is generated, a DataFrame is returned
```

```
grouped.mean()
```

```
Out[23]:
```

		total_bill	tip	size
sex	smoker			
Female	No	18.105185	2.773519	2.592593
	Yes	17.977879	2.931515	2.242424
Male	No	19.791237	3.113402	2.711340
	Yes	22.284500	3.051167	2.500000

```
In [24]: # The returned results can be further manipulated via apply()
# For example, suppose the bill and tips are in USD but we want EUR
```

```
import json
import urllib
```

```
# get current conversion rate
```

```
converter = json.loads(urllib.urlopen('http://rate-exchange.appspot.com/currency?from=USD&to=EUR').read())
```

```
print converter
```

```
grouped['total_bill', 'tip'].mean().apply(lambda x: x*converter['rate'])
```

```
{u'to': u'EUR', u'rate': 0.879191, u'from': u'USD'}
```

```
Out[24]:
```

		total_bill	tip
sex	smoker		
Female	No	15.917916	2.438453
	Yes	15.805989	2.577362
Male	No	17.400278	2.737275
	Yes	19.592332	2.682558

```
In [25]: # We can also transform the original data for more convenient analysis
# For example, suppose we want standardized units for total bill and tips
```

```
zscore = lambda x: (x - x.mean())/x.std()
```

```
std_grouped = grouped['total_bill', 'tip'].transform(zscore)
```

```
std_grouped.head(n=4)
```

```
Out[25]:
```

	total_bill	tip
0	-0.153049	-1.562813
1	-1.083042	-0.975727
2	0.139661	0.259539
3	0.445623	0.131984

```
In [26]: # Suppose we want to apply a set of functions to only some columns
```

```
grouped['total_bill', 'tip'].agg(['mean', 'min', 'max'])
```

```
Out[26]:
```

		total_bill			tip		
		mean	min	max	mean	min	max
sex	smoker						
Female	No	18.105185	7.25	35.83	2.773519	1.00	5.2
	Yes	17.977879	3.07	44.30	2.931515	1.00	6.5
Male	No	19.791237	7.51	48.33	3.113402	1.25	9.0
	Yes	22.284500	7.25	50.81	3.051167	1.00	10.0

```
In [27]: # We can also apply specific functions to specific columns
df = grouped.agg({'total_bill': (min, max), 'tip': sum})
df
```

```
Out[27]:
```

		tip	total_bill	
		sum	min	max
sex	smoker			
Female	No	149.77	7.25	35.83
	Yes	96.74	3.07	44.30
Male	No	302.00	7.51	48.33
	Yes	183.07	7.25	50.81

0.1.5 Using statsmodels

Many of the basic statistical tools available in R are replicated in the `statsmodels` package. We will only show one example.

```
In [28]: # Simulate the genotype for 4 SNPs in a case-control study using an additive genetic model

n = 1000
status = np.random.choice([0,1], n )
genotype = np.random.choice([0,1,2], (n,4))
genotype[status==0] = np.random.choice([0,1,2], (sum(status==0), 4), p=[0.33, 0.33, 0.34])
genotype[status==1] = np.random.choice([0,1,2], (sum(status==1), 4), p=[0.2, 0.3, 0.5])
df = DataFrame(np.hstack([status[:, np.newaxis], genotype]), columns=['status', 'SNP1', 'SNP2', 'SNP3', 'SNP4'])
df.head(6)
```

```
Out[28]:
```

	status	SNP1	SNP2	SNP3	SNP4
0	0	2	1	2	0
1	1	1	0	2	2
2	1	0	1	2	1
3	1	2	2	1	2
4	1	1	2	0	1
5	1	0	0	1	2

```
In [29]: # Use statsmodels to fit a logistic regression to the data
fit1 = sm.Logit.from_formula('status ~ %s' % '+'.join(df.columns[1:]), data=df).fit()
fit1.summary()
```

```
Optimization terminated successfully.
Current function value: 0.642824
Iterations 5
```

```
Out[29]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        Logit Regression Results
=====
Dep. Variable:                  status   No. Observations:                  1000
```



```

Model:                      Logit      Df Residuals:          995
Method:                     MLE        Df Model:              4
Date:                       Thu, 22 Jan 2015    Pseudo R-squ.:        0.07259
Time:                       15:34:43    Log-Likelihood:       -642.82
converged:                   True        LL-Null:              -693.14
                                   LLR p-value:      7.222e-21
=====
              coef      std err          z      P>|z|      [95.0% Conf. Int.]
-----
Intercept    -1.7409      0.203     -8.560      0.000      -2.140      -1.342
SNP1          0.4306      0.083      5.173      0.000       0.267       0.594
SNP2          0.3155      0.081      3.882      0.000       0.156       0.475
SNP3          0.2255      0.082      2.750      0.006       0.065       0.386
SNP4          0.5341      0.083      6.404      0.000       0.371       0.698
=====
"""

```

```

In [30]: # Alternative using GLM - similar to R
fit2 = sm.GLM.from_formula('status ~ SNP1 + SNP2 + SNP3 + SNP4', data=df, family=sm.families.B)
print fit2.summary()
print chisqprob(fit2.null_deviance - fit2.deviance, fit2.df_model)
print(fit2.null_deviance - fit2.deviance, fit2.df_model)

```

Generalized Linear Model Regression Results

```

=====
Dep. Variable:              status    No. Observations:          1000
Model:                     GLM        Df Residuals:            995
Model Family:              Binomial   Df Model:                4
Link Function:              logit     Scale:                  1.0
Method:                     IRLS      Log-Likelihood:         -642.82
Date:                       Thu, 22 Jan 2015    Deviance:              1285.6
Time:                       15:34:43    Pearson chi2:          1.01e+03
No. Iterations:              5
=====
              coef      std err          t      P>|t|      [95.0% Conf. Int.]
-----
Intercept    -1.7409      0.203     -8.560      0.000      -2.140      -1.342
SNP1          0.4306      0.083      5.173      0.000       0.267       0.594
SNP2          0.3155      0.081      3.882      0.000       0.156       0.475
SNP3          0.2255      0.082      2.750      0.006       0.065       0.386
SNP4          0.5341      0.083      6.404      0.000       0.371       0.698
=====
7.22229516479e-21
(100.63019840179481, 4)

```

0.2 Using R from IPython

While Python support for statistical computing is rapidly improving (especially with the pandas, statsmodels and scikit-learn modules), the R ecosystem is still vastly larger. However, we can have our cake and eat it too, since IPython allows us to run R (almost) seamlessly with the Rmagic (rpy2.ipython) extension.

There are two ways to use Rmagic - using %R (applies to single line) and %%R (applies to entire cell). Python objects can be passed into R with the -i flag and R objects passed out with the -o flag.

```

In [31]: ! pip install ggplot &> /dev/null

```

0.2.1 Using Rmagic

```
In [32]: %load_ext rpy2.ipython
```

```
In [33]: %%R -i df,status -o fit
```

```
fit <- glm(status ~ ., data=df)
print(summary(fit))
print(fit$null.deviance - fit$deviance)
print(fit$df.null - fit$df.residual)
with(fit, pchisq(null.deviance - deviance, df.null - df.residual, lower.tail = FALSE))
```

Call:

```
glm(formula = status ~ ., data = df)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-0.7927	-0.4464	0.2073	0.4301	0.8999

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.10014	0.04323	2.316	0.02075 *
SNP1	0.09904	0.01874	5.285	1.55e-07 ***
SNP2	0.07217	0.01836	3.932	9.01e-05 ***
SNP3	0.05135	0.01856	2.767	0.00576 **
SNP4	0.12372	0.01869	6.620	5.86e-11 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.2269642)

Null deviance: 250.00 on 999 degrees of freedom
Residual deviance: 225.83 on 995 degrees of freedom
AIC: 1361.9

Number of Fisher Scoring iterations: 2

```
[1] 24.16657
[1] 4
[1] 7.396261e-05
```

Using rpy2 directly

```
In [34]: import rpy2.robjects as ro
        from rpy2.robjects.packages import importr
```

```
base = importr('base')
```

```
fit_full = ro.r("lm('mpg ~ wt + cyl', data=mtcars)")
print(base.summary(fit_full))
```

Call:

```
lm(formula = "mpg ~ wt + cyl", data = mtcars)
```

```

Residuals:
    Min       1Q   Median       3Q      Max
-4.2893 -1.5512 -0.4684  1.5743  6.1004

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  39.6863     1.7150  23.141 < 2e-16 ***
wt          -3.1910     0.7569  -4.216 0.000222 ***
cyl         -1.5078     0.4147  -3.636 0.001064 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.568 on 29 degrees of freedom
Multiple R-squared:  0.8302,    Adjusted R-squared:  0.8185
F-statistic: 70.91 on 2 and 29 DF,  p-value: 6.809e-12

```

0.2.2 Using R from pandas

Reading R dataset into Python

```
In [35]: import pandas.rpy.common as com
```

```

df = com.load_data('mtcars')
print df.head(n=6)

```

```

mpg  cyl  disp  hp  drat    wt    qsec  vs  am  gear  carb
0  21.0    6   160  110   3.90  2.620  16.46   0   1     4     4
1  21.0    6   160  110   3.90  2.875  17.02   0   1     4     4
2  22.8    4   108   93   3.85  2.320  18.61   1   1     4     1
3  21.4    6   258  110   3.08  3.215  19.44   1   0     3     1
4  18.7    8   360  175   3.15  3.440  17.02   0   0     3     2
5  18.1    6   225  105   2.76  3.460  20.22   1   0     3     1

```

```
In [36]: %load_ext version_information
```

```
%version_information numpy, matplotlib, pandas, statsmodels
```

Out[36]:

Software	Version
Python	2.7.9 64bit [GCC 4.2.1 (Apple Inc. build 5577)]
IPython	2.3.1
OS	Darwin 13.4.0 x86_64 i386 64bit
numpy	1.9.1
matplotlib	1.4.2
pandas	0.15.1
statsmodels	0.5.0
Thu Jan 22 15:34:45 2015 EST	